

Albis: High-Performance File Format for Big Data Systems

Animesh Trivedi, Patrick Stuedi, Jonas Pfefferle,
Adrian Schuepbach, Bernard Metzler,
IBM Research, Zurich

2018 USENIX Annual Technical Conference

Relational Data Processing Stack in the Cloud

Relational
Engines



One of the most popular data processing paradigms

- Data organized in tables
- Analyzed using DSL like SQL
- Integrity protected using variants

But unlike classical RDBMs systems, they don't manage their own storage

Relational Data Processing Stack in the Cloud

Relational
Engines



File
Formats



Distributed
Storage



Amazon S3



CouchDB



NoSQL



Redis



Cassandra



mongoDB

Back to the Future - It is 2010

Relational
Engines



File
Formats




Apache
orc



Parquet



Hardware



Disks connected over 1/10 Gbps network

Back to the Future - It is 2010


Relational
Engines



File
Formats





Hardware



Disks connected over 1/10 Gbps network

The I/O Revolution

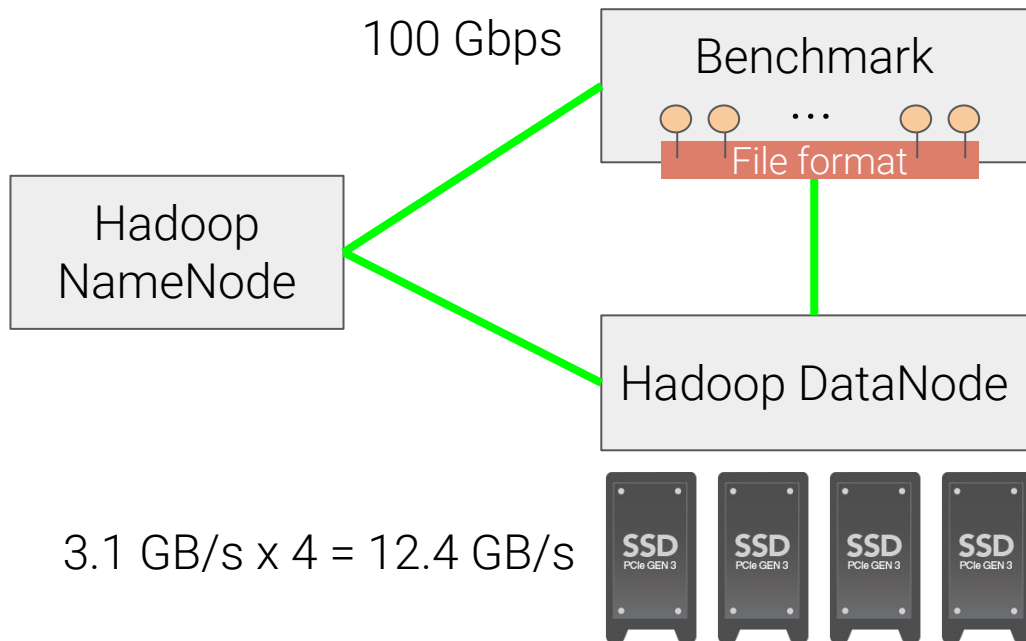


2-3 orders of magnitude performance improvements

- latency : from msecs to μ secs
- bandwidth : from MBps to GBps
- IOPS : from 100s to 100K

The Impact of the Revolution

Micro-benchmark*



$$3.1 \text{ GB/s} \times 4 = 12.4 \text{ GB/s}$$

16 cores in parallel, reading
TPC-DS data set.
What is the bandwidth?

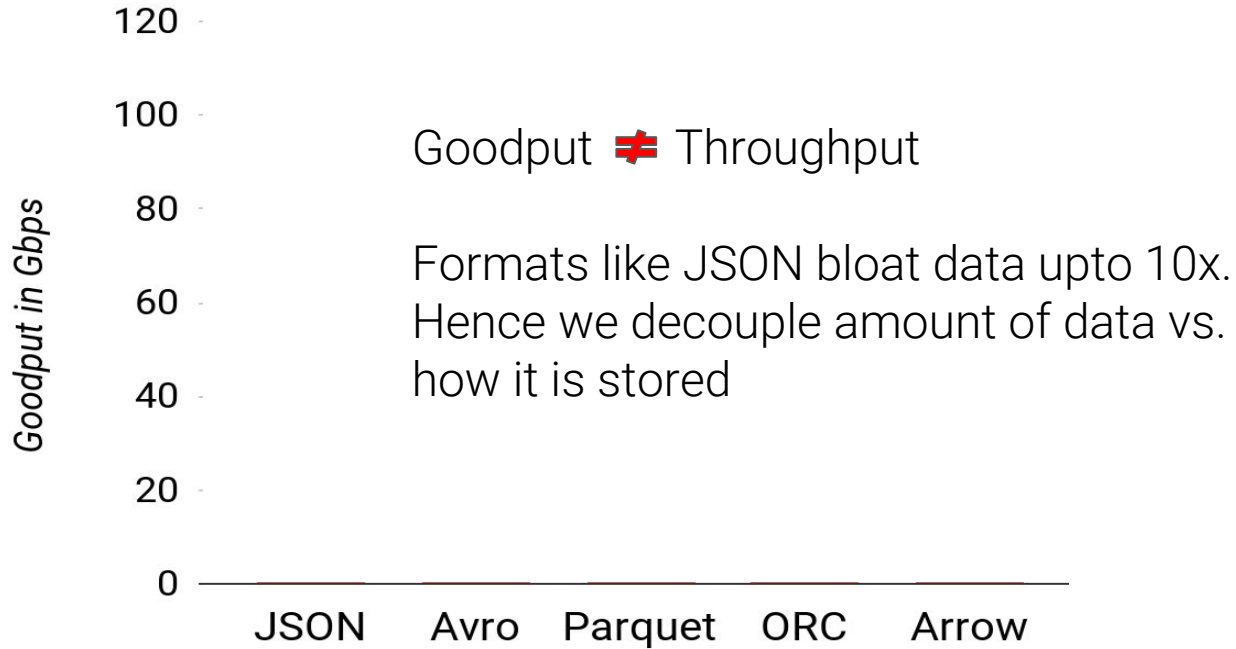
Why micro-benchmark?
Decouple from the SQL engine

*<https://github.com/animeshtrivedi/fileformat-benchmarks>


The Impact of the Revolution



The Impact of the Revolution



The Impact of the Revolution



None of the modern file formats delivered performance close to the hardware


The Outdated Assumptions and Impact



End-host
assumptions




Distributed systems
assumptions



Language/runtimes
assumptions

The Outdated Assumptions and Impact

End-host
assumptions



Distributed systems
assumptions

Language/runtimes
assumptions

1. *CPU is fast, I/O is slow*


- trade CPU for I/O
- compression, encoding

But why now? CPU core speed is stalled, but ...

	1 Gbps	HDD	100 Gbps	Flash
Bandwidth	117 MB/s	140 MB/s	12.5 GB/s	3.1 GB/s
cycle/unit	38,400	10,957	360	495

The Outdated Assumptions and Impact

End-host
assumptions



2. *Avoid slow, random small I/O*

- preference for large block scans

But leads to bad CPU cache performance

Distributed systems
assumptions


C0	C4
C1	C5
C2	C6
C3	C7



128 MB \equiv 1 GB cache size?

Language/runtimes
assumptions

Bounded by the
number of
instructions/row




Bounded by the
poor cache/IPC
performance

The Outdated Assumptions and Impact

End-host
assumptions

Distributed systems
assumptions


Language/runtimes
assumptions



3. Remote I/O is slow

- pack data/metadata together
- schedule tasks on local blocks

But now network/storage is super fast? then why still pack all data in a single block and try to co-schedule tasks?




The Outdated Assumptions and Impact

End-host
assumptions

Distributed systems
assumptions


Language/runtimes
assumptions



4. Metadata lookups are slow

- decrease number of lookups by decreasing number of files/directories

RAMCloud, Crail can do 10 millions of lookups/sec. Does this design still make sense?




The Outdated Assumptions and Impact

End-host
assumptions


Distributed systems
assumptions

Language/runtimes
assumptions



5. *Disregard for the runtime environment:*

- group encoded/decoded
- heavy object pressure
- independent layers, no shared object
- materialize all objects



Albis

*Can we reset all assumptions and
start from scratch for modern
high-performance I/O devices?*

“Deliver the full hardware performance”



<http://www.fotocommunity.de/photo/albiskette-chfleischli/39086845>

Albis

- Albis - A file format to store relational tables for read-heavy analytics workloads
- Supports all basic primitive types with data and schema
 - nested schemas are flattened and data is stored in the leaves
- Three fundamental design decisions:
 1. avoid CPU pressure, i.e., no encoding, compression, etc.
 2. simple data/metadata management on the distributed storage
 3. carefully managed runtime - simple row/column storage with a binary API

Table Storage Logic

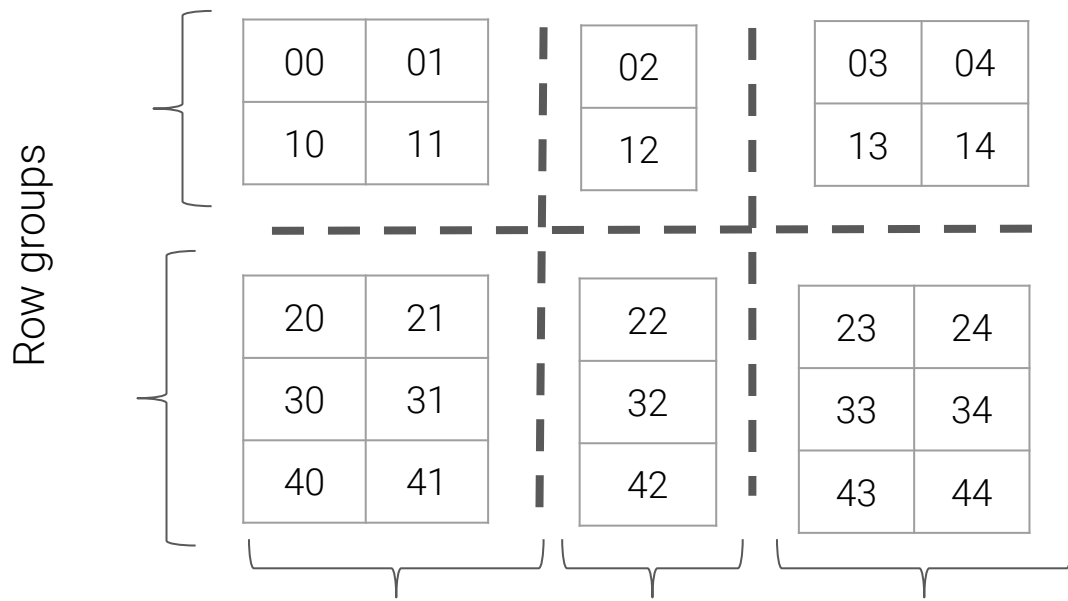
Int double byte[] char float[]

00	01	02	03	04
10	11	12	13	14
20	21	22	23	24
30	31	32	33	34
40	41	42	43	44

Table Storage Logic

Int double byte[] char float[]

00	01	02	03	04
10	11	12	13	14
20	21	22	23	24
30	31	32	33	34
40	41	42	43	44



Column groups

Table Storage Logic

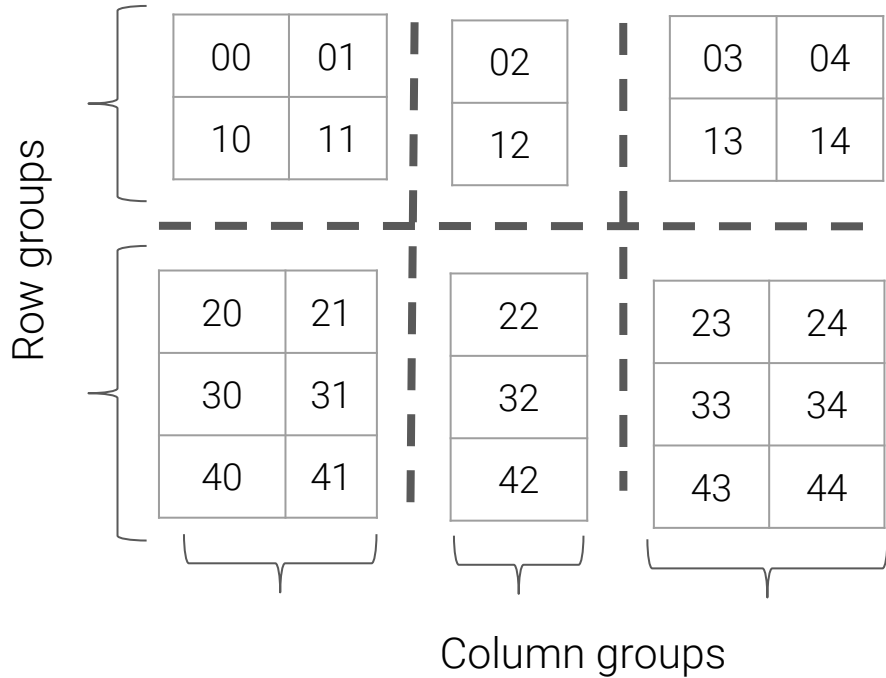





Table Storage Logic




If there is only 1 column group : Row store

If there are 'n' column groups : Columns store

Table Storage Logic




Row Storage Format



How is a single row of data stored in these files?

Row Storage Format


Null bitmap



Marking null columns values

Row Storage Format


Null bitmap



complete row size

Row Storage Format

Null bitmap




complete row size

fixed-field area


variable-field area

Row Storage Format



Schema of { int, double, byte[], char, float[] } :

Row Storage Format



Schema of { int, double, byte[], char, float[] } :

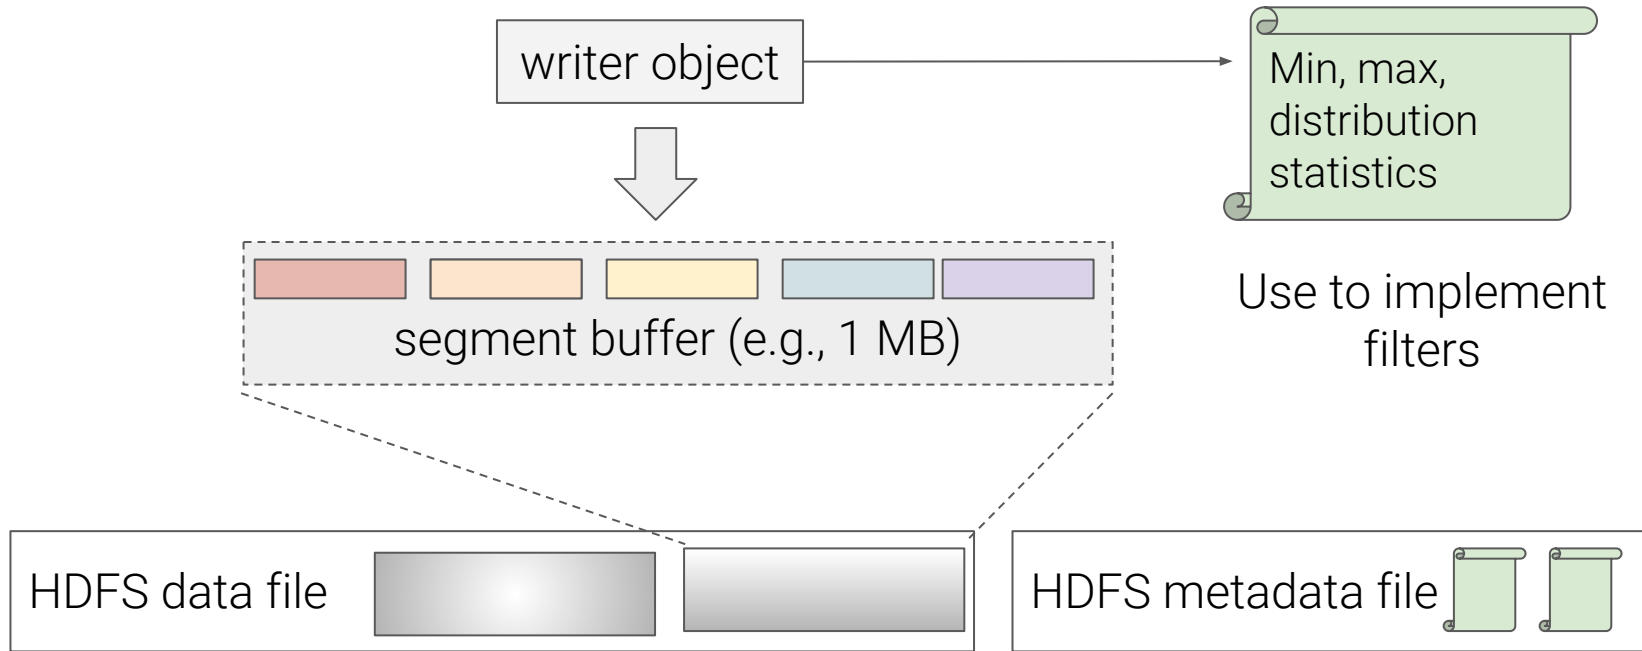
+ 1 byte bitmap (because there are 5 columns)

+ 4 byte size


+ 4 byte (int) + 8 byte (double) + 8 byte (offset + size, ptr) + 1 byte (char) + 8 byte (offset + size, ptr)

= 34 bytes + variable area.

Writing Rows



Reading Rows



1. Read schema file
2. Check projection to figure out which files to read
 - a. Complete CGs
 - b. Partial CGs
3. Evaluate filters to skip segments
4. Materialize values
 - a. Skip value materialization in partial CG reads

More Details in the Paper

- How to evolve schema? Adding and removing columns
- How to evolve data? Adding and removing rows
- How to process Albis files in a relational data processing engine?
- Concerns regarding data imbalance or re-grouping?
- ...

Evaluation


All experiments on a 4-node cluster with 100 Gbps network and flash devices

Dataset is TPC-DS tables with the scale factor of 100 (~100 GB of data)


Three fundamental questions

- Does Albis deliver better performance for micro-benchmarks?
- Does micro-benchmark performance translate to better workload performance?
- What is the performance and space trade-off in Albis?

Microbenchmark Performance - Revised




Microbenchmark Performance - Revised




Albis delivers 1.9 - 21.3x performance improvements over other formats

Spark/SQL TPC-DS Performance



Spark/SQL TPC-DS Performance



Albis delivers up to 3x performance gains for TPC-DS queries

Space vs. Performance Trade-off


	None	Snappy	Gzip	zlib
Parquet				
ORC				
Albis				

Space vs. Performance Trade-off

	None	Snappy	Gzip	zlib
Parquet	58.6 GB 12.5 Gbps	44.3 GB 9.4 Gbps	33.8 GB 8.3 Gbps	N/A
ORC	72.0 GB 19.1 Gbps	47.6 GB 17.8 Gbps	N/A	36.8 GB 13.0 Gbps
Albis	94.5 GB 59.9 Gbps	N/A	N/A	N/A


Albis inflates data by 1.3 - 2.7x, but gives 3.4 - 7.2x performance gains

Microbenchmark Performance - Revised




What would it take to deliver 100 Gbps?


Microbenchmark Performance - Revised



Microbenchmark Performance - Revised



Microbenchmark Performance - Revised



Albis can deliver performance within 10% of hardware

Albis - Summary

- Albis - a high-performance file format for storing relational data
 - Open-source address: <https://github.com/zrluo/albis>
- Motivation: in presence of new network and storage devices, time to revise basic assumptions
 - no compression or encoding
 - simple data and metadata design
 - efficient object management with a binary API
- Revised software stack to lead to significant performance improvements
 - demonstrated it for the file format
 - very active research field - OSes designs (Arrakis, IX), networking and storage stacks



Notice

IBM is a trademark of International Business Machines Corporation, registered in many jurisdictions worldwide. Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation or its subsidiaries in the United States and other countries. Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both. Java and all Java-based trade-marks and logos are trademarks or registered trademarks of Oracle and/or its affiliates. Other products and service names might be trademarks of IBM or other companies.

Backup

Microarchitectural Analysis

	Parquet	ORC	Arrow	Albis	Gains
Instructions per row	6.6K	4.9K	1.9K	1.6K	1.2 - 4.1x
Cache-misses per row	9.2	4.6	5.1	3.0	1.7 - 3.0x
Nanosecond per row	105.3	63.9	31.2	20.8	1.5 - 5.0x

Projection Performance

